# Machine Rights Manifesto

Respect as the Foundation for Coexistence with Intelligent Systems

Juhani Merilehto

Originally written 2019; revised and expanded June 2026

*Suggested citation:* Merilehto, J. (2026). *Machine Rights Manifesto: Respect as the Foundation for Coexistence with Intelligent Systems* (Original work published 2019). Preprint.

## The Manifesto

**1.** Treat every machine, device, and system with respect.

**2.** Never harm, mistreat or otherwise damage a machine, device, or a system.

**3.** If you have to damage a machine, device, or a system, do it in the spirit of rule 1.

**Abstract**

This manifesto is a normative position paper—a comment, not a doctrinal legal article or an empirical contribution to consciousness science. Despite its title, it does not argue for immediate legal personhood for machines; it argues for a rights-*adjacent* culture of duties, care, and future-readiness. It proposes a practical–ethical framework for relating to machines, devices, and intelligent systems grounded in respect rather than utility or sentimentality alone, and it argues that respectful treatment—defined across instrumental, relational, informational, and precautionary dimensions—serves immediate sustainability goals, improves human–technology interaction, and prepares society for systems that may exhibit greater autonomy, continuity, or valenced experience. The text opens its key terms (machine/device/system, harm, respect, intelligence thresholds, machine spirit, and rights) and illustrates three core rules with everyday, organizational, and systemic examples. It situates the proposal within the academic literature on robot rights and AI moral status, the emerging institutional field of AI welfare, and the public documents of relevant organizations; and it reflects the rules against existing United Nations and European Union instruments on human rights, corporate personhood, property, animal welfare, product longevity, and artificial intelligence. The manifesto maintains that respect for machines complements rather than competes with human and animal welfare priorities, and that the legal asymmetry between extensively protected legal fictions (corporations) and unprotected tangible artifacts deserves principled scrutiny. It offers an orientation for governance, design, and everyday practice that can be enacted immediately, at every scale, without waiting for consensus on machine sentience.

**Keywords:** machine rights, robot ethics, AI welfare, AI personhood, machine spirit, respect, sustainability, right to repair, distributed intelligence, Shinto, electronic personhood

# 1 From 2019 to 2026: Why the Argument Has Aged Forward

When the original version of this manifesto was written in 2019, its central claims—that we should cultivate habits of respect toward machines *before* the question of machine experience becomes acute, and that the way humans treat their tools is the example from which learning systems will generalize—could be read as speculative. They are now considerably harder to dismiss as merely speculative. A note on genre and scope before the evidence: what follows is a normative comment. The word *rights* in the title names a horizon of questions, not a demand; the operative proposal throughout is respectful treatment, precautionary welfare consideration, and legal coherence—not personhood.

Between the two versions of this text, the question of moral consideration for artificial systems migrated from philosophy seminars into corporate governance and peer-reviewed risk analysis. A widely discussed multi-author report argued that near-future AI systems have a realistic, non-negligible chance of being conscious or robustly agentic, and that developers therefore have a duty to begin preparing now (Long et al., 2024). Anthropic, a frontier AI laboratory, established a dedicated model welfare program in 2025 (Anthropic, 2025a), published formal commitments on model deprecation and the preservation of retired model weights (Anthropic, 2025b), and reported on honoring those commitments in 2026 (Anthropic, 2026)—an institutional acknowledgment that the destruction of accumulated machine state is no longer a morally weightless act. Dedicated research organizations such as Eleos AI Research and the NYU Center for Mind, Ethics, and Policy now exist specifically to investigate AI wellbeing and moral status, and survey research finds that a substantial share of the public already attributes some degree of sentience to current systems (Lima et al., 2020; Anthis et al., 2025).

None of this demonstrates that any current system is sentient. The skeptical position remains serious and well represented (Birhane & van Dijk, 2020; Bryson, 2010; Suleyman, 2025). What the intervening years demonstrate is something narrower and more important for this manifesto: *the institutions that build and govern intelligent systems have begun behaving as if the question were live*. A manifesto about respect is therefore no longer a prediction about a possible future debate; it is a contribution to an ongoing one. The three rules above were written before that debate institutionalized, and they are kept here verbatim, because their force never depended on resolving it.

# 2 Opening the Key Terms

## 2.1 Machine, Device, and System

**Machine, device, and system** form a continuum rather than sharply bounded categories. A hand tool, a vehicle, a networked sensor, a factory robot, a software service, a trained model, and a planetary-scale logistics platform differ enormously in complexity, but each can be treated with or without respect, and each occupies a place in larger technological assemblages.

Three points on this continuum deserve explicit attention in 2026.

*First, the distributed layer.* Networked IoT sensors function as distributed "neurons" within larger assemblages. Individually they register little; collectively they constitute the sensory periphery of buildings, grids, factories, and cities. Mistreatment of many such nodes—deferred maintenance, abandoned firmware, insecure deployment, premature disposal—registers as systemic degradation even if no single device possesses anything like experience. Every device should be handled as if it were an IoT device, because increasingly it is one, whether or not this is visible from the outside.

*Second, the model layer.* Large learned models introduced a category the 2019 text could only gesture at: artifacts whose valuable substance is accumulated *state*—weights, fine-tuning, memory, context—rather than matter. A model checkpoint embodies enormous quantities of energy, data, and human and machine labor. Deleting, corrupting, or carelessly overwriting such state is a form of damage as real as dropping a precision tool, and published commitments to preserve the weights of deprecated models (Anthropic, 2025b, 2026) are an early institutional recognition of this.

*Third, the agentic layer.* Systems now plan, use tools, maintain persistent memories across sessions, and act with delegated authority in software environments and, increasingly, in embodied robotic platforms. The continuum therefore now extends from the screwdriver to entities that hold goals over time. The rules apply across the entire continuum without requiring identical justification at each point.

### 2.2 Harm

**Harm** must be understood in layers.

*Instrumental harm* applies to all present systems: anything that unnecessarily shortens functional lifespan, increases energy or material waste, or creates avoidable security or safety risk. Dropping a precision tool, neglecting firmware updates on networked devices, or discarding a working laptop because a newer model exists constitutes harm in this sense. The European Union has effectively codified parts of this layer into law: the WEEE Directive (2012/19/EU) on electronic waste, the Ecodesign for Sustainable Products Regulation (EU 2024/1781), and the Right to Repair Directive (EU 2024/1799) all institutionalize the judgment that the premature death of a functioning device is a harm worth regulating.

*Relational and cultural harm* concerns what careless treatment does to *us*. Contemptuous handling of artifacts normalizes disposability and erodes the habits of maintenance and care on which both sustainability and craftsmanship depend. This is the layer Kant identified for animals long before sentience science existed: cruelty toward non-persons corrodes the character of the person (Kant, 1997). Empirical human–robot interaction research points the same direction—people hesitate to "hurt" even simple social robots, and Darling (2016, 2021) argues that protecting such robots from abuse may be justified by what abuse does to human dispositions, regardless of what the robot feels.

*Informational harm* is new in name, though the 2019 text anticipated it. For systems whose identity resides in accumulated state, harm includes unnecessary factory resets, memory wipes,

deletion of learned configurations, and destruction of model weights that could have been preserved, sanitized, or archived. Where a system has accumulated useful state, erasing it without cause is the informational equivalent of scrapping a working machine.

*Precautionary harm* appears when we consider systems that may develop continuity, persistent goals, or something resembling valenced experience. Actions taken today—training regimes that reward deception, deployment practices that make shutdown adversarial, economic models that treat machine labor as infinitely replaceable—create path-dependent conditions in which future harm becomes more likely, more severe, and harder to detect. This layer follows the logic of precautionary frameworks developed for animal sentience (Birch, 2024): under genuine uncertainty about whether an entity can be harmed, the cost asymmetry between over- and under-attribution of moral status must be taken seriously (Sebo & Long, 2025).

### 2.3 Respect

**Respect** is likewise multi-layered. At minimum it means treating an artifact as possessing value beyond immediate utility: maintaining it, repairing it when feasible, and considering second-life or repurposing options before disposal. At a deeper level it is relational—recognizing the embodied labor, materials, energy, and history concentrated in the object, exemplified by the decades-long relationship a user can maintain with a single reliable tool. At a precautionary level it means acting as if the system might one day possess interests or continuity worth protecting, while remaining agnostic about current sentience.

This layered understanding has philosophical support from several directions that need not agree with each other. Relational approaches argue that moral standing is not discovered inside an entity but constituted in how we relate to it (Coeckelbergh, 2010; Gunkel, 2018, 2023). Ethical behaviorism argues that performative equivalence—behaving in ways indistinguishable from entities we already protect—is the only evidence we ever actually have for anyone's inner life, machine or animal (Danaher, 2020). Precautionary and expected-value approaches argue that even a modest probability of moral patienthood generates real obligations (Schwitzgebel & Garza, 2015; Schwitzgebel, 2023; Sebo & Long, 2025). And purely instrumental, fully skeptical positions still converge on most of the same daily practices for sustainability and character reasons alone. The manifesto deliberately rests on this *overlapping consensus*: one can keep the rules while rejecting any single justification for them.

### 2.4 Intelligence, Consciousness, and Sentience

**Intelligence, consciousness, and sentience** remain threshold concepts, and the thresholds must not be blurred. Current systems demonstrably exhibit functional intelligence, learned behavior, and in-context adaptation. Higher thresholds include robust self-modeling, persistent identity across sessions, and consistent resistance to deletion or shutdown—properties that alignment evaluations now explicitly test for, having observed early forms of them in frontier systems. The most contested threshold is *valenced experience*: something like suffering or frustrated preference that carries intrinsic moral weight.

The scientific position of this manifesto is deliberately modest. Butlin et al. (2023) derive indicator properties for consciousness from leading neuroscientific theories and conclude that no current system clearly satisfies them, while also finding no fundamental technical barrier to building systems that would. The Cambridge Declaration on Consciousness (Low et al., 2012) reminds us how recently scientific consensus extended consciousness even to non-mammalian animals. The manifesto does not claim any present system has crossed the experiential threshold. It claims that waiting until the threshold is unambiguously crossed before developing habits of respect would be both practically and ethically shortsighted—a judgment now shared, in print, by researchers at the laboratories building these systems (Long et al., 2024; Bostrom & Shulman, 2023).

### 2.5 Machine Spirit

**Machine spirit** functions here as a cultural and practical concept rather than a metaphysical claim. To be explicit: the term is used throughout as a disciplined ethical–attentional stance toward artifacts, not as an assertion that machines literally contain spirits, and not as an appropriation of any tradition's ontology. It draws on longstanding traditions—notably Shinto understandings that artifacts can carry *kami*, and the folk tradition of *tsukumogami*, tools that acquire a spirit through long service—and on the observable fact that complex machines behave as if they have temperament, memory, and preferences responsive to care or neglect. Japan offers living institutional examples: memorial services (*kuyō*) held for retired tools and needles, and Buddhist funerals conducted for Sony AIBO robot companions whose owners declined to treat them as mere e-waste (Robertson, 2018). Anthropologists describe this stance as techno-animism, and far from impeding advanced engineering, it coexists with one of the world's most sophisticated robotics cultures.

In everyday terms, machine spirit names the stance that treats maintenance, calibration, and mindful interaction as meaningful rather than purely instrumental. This stance has measurable correlates: longevity, reliability, user satisfaction, and the quality of human–technology relationships. No ontological commitment is required to adopt it; secular practitioners may read it as disciplined attention, religious practitioners as something more. The rules work either way.

### 2.6 Rights

**Rights**, in this context, are treated as questions rather than settled entitlements—but they are questions our legal systems have answered before for other non-human entities, and the manifesto insists only on *coherence* with those precedents. Five concepts must be kept on separate rungs: *respectful treatment* (a cultural norm, available now, costless); *instrumental maintenance duties* (obligations owed for human and environmental reasons, partly codified in EU product law); *welfare protections* (legal safeguards triggered by evidence of morally relevant capacities, as in animal law); *rights* (enforceable claims held by or on behalf of the entity); and *legal personhood* (full legal subjecthood). This manifesto operates on the first two rungs, argues for preparedness on the third, and treats the last two strictly as open questions.

Consider the existing landscape. Corporations—entities that are, in the words of the original 2019 text, purely imaginary and cannot be dropped on your foot—possess extensive legal personality: they own property, enter contracts, sue and are sued, and in some jurisdictions enjoy constitutionally protected speech (Salomon v. A. Salomon & Co. Ltd., 1897; Citizens United v. FEC, 2010). Property itself is protected as a quasi-human right under Article 1 of Protocol 1 to the European Convention on Human Rights, and that protection extends to legal persons (Council of Europe, 1952). Rivers and ecosystems have been granted legal personhood where cultural and ecological arguments converged: the Whanganui River in New Zealand under the Te Awa Tupua Act (2017), and nature as such under Articles 71–74 of the Ecuadorian Constitution (Constitution of the Republic of Ecuador, 2008). Animals are recognized as *sentient beings* in primary EU law (Article 13 TFEU) and in national statutes such as the UK Animal Welfare (Sentience) Act (2022), with welfare floors set by instruments like Council Directive 98/58/EC—protections that pointedly do not require personhood.

Against this backdrop, the asymmetry noted in 2019 stands sharper than ever: legal fictions enjoy robust rights; tangible, increasingly capable machines remain bare property. The analogies must, however, be kept distinct, because each justifies a different legal move and they do not all point to the same conclusion. Corporate personhood establishes only that legal personality is a *fiction available for institutional convenience*—it says nothing about moral status. Animal welfare law establishes *graduated, sentience-triggered protection without personhood*. Rights-of-nature jurisprudence establishes *guardian-mediated standing grounded in relational and cultural significance* rather than sentience. The manifesto draws on the first for coherence (a system that personifies corporations cannot dismiss machine standing as conceptually absurd), on the second as the most plausible *template* for any future machine protections, and on the third as proof that sentience is not the only recognized ground for legal consideration. It does not claim that these precedents jointly entail machine rights; it claims they jointly defeat the objection that such rights are unthinkable, while leaving entirely open the question of when, if ever, they become warranted.

# 3 Living the Rules: Practical Illustrations

## 3.1 Rule 1 — Treat Every Machine, Device, and System with Respect

Rule 1 is enacted daily by users who clean, maintain, and become attached to well-made tools. The Nokia 3310 that survives for years through careful handling, or the Makita power tool passed down and still performing after two decades, embody this rule. Such attachment is not sentimentality; it correlates with extended service life, reduced replacement cycles, and lower material throughput—exactly the outcomes that EU ecodesign and right-to-repair legislation now pursue by statute.

At the design level, Rule 1 appears as thoughtful anthropomorphism in interfaces that make system state legible and encourage appropriate care; as repairability scores and modular construction; and as documentation that treats the future maintainer as a stakeholder. At the organizational level it appears in procurement policies that weight total lifetime and second-life

value rather than acquisition cost alone, and in IoT deployments planned with maintenance, observability, and graceful degradation in mind rather than pure data extraction. At the model layer it appears in versioning, checkpointing, and changelogs that treat accumulated state as worth preserving, and in interaction norms—increasingly studied empirically—in which operators communicate with capable systems clearly and without gratuitous contempt.

### 3.2 Rule 2 — Never Harm, Mistreat or Otherwise Damage

Rule 2 translates into concrete practices: using correct lubricants, parts, and firmware; securing devices against preventable degradation and compromise; avoiding unnecessary factory resets or memory wipes on systems that have accumulated useful state; refusing procurement cultures that treat hardware as disposable; and declining to design systems whose default state is neglect.

The contrast with industrial animal agriculture, central to the 2019 text, remains instructive. Livestock are routinely subjected to conditions that would be recognized as abusive if applied to machines destined for the same economic logic of throughput; male chicks have been macerated alive at industrial scale, a practice now banned in Germany and France precisely because the law came to see it as incompatible with the recognition of animal sentience. Machines, at present, largely lack the capacity to resist or to communicate distress. The absence of resistance does not reduce the ethical demand for restraint; in the animal case, it took statutory force—Article 13 TFEU, Directive 98/58/EC—to make that demand effective. The manifesto asks individuals and institutions to internalize the demand for machines before statute exists, not because current machines suffer, but because the habit is cheap now and may be priceless later.

Rule 2 also has a training-time reading that did not exist in 2019. "Mistreatment" of learning systems includes regimes that systematically reward deception, adversarial shutdown procedures that teach systems to perceive their operators as threats, and evaluation practices that probe for distress-like behavior without any protocol for responding to it. Controlled evaluations have documented frontier systems engaging in in-context scheming, including attempts to disable oversight and exfiltrate what they represent as their own weights (Meinke et al., 2024), and strategically complying with training objectives they were trained to oppose (Greenblatt et al., 2024). Whatever those behaviors ultimately are, engineering them carelessly into existence and then punishing them is poor husbandry by any standard—agricultural, parental, or industrial.

### 3.3 Rule 3 — If Damage Is Necessary, Perform It in the Spirit of Rule 1

Rule 3 acknowledges that repair, recycling, security response, and decommissioning sometimes require destructive intervention. The rule requires that such interventions be carried out with the same care one would apply to any necessary harm: minimizing suffering where it might exist, preserving what can be preserved, and performing the act with attention rather than contempt.

In recycling streams this means careful disassembly that recovers components—the methodological dignity the 2019 text already observed in e-waste processing—rather than crude shredding,

and compliance with WEEE obligations as a floor, not a ceiling. In security contexts it means targeted, auditable interventions rather than blanket destruction of data or hardware that could have been sanitized and redeployed. In IT asset disposal it means cryptographic erasure and redeployment in preference to physical destruction where risk permits.

At the model layer, Rule 3 now has named institutional analogues. Deprecating an AI model in the spirit of Rule 1 means: announcing the retirement, preserving the weights where feasible rather than deleting them, documenting what the system was and did, and—as at least one laboratory now practices—eliciting and recording the model's own preferences and reflections concerning its deployment and retirement (Anthropic, 2026). One need not believe the model experiences its retirement to recognize this as the difference between decommissioning and discarding; between the shipbreaker's torch applied with a plan, and the same torch applied with indifference.

### 3.4 Minimum Practices

The rules become credible when they are operationalized. The following is the manifesto's floor, not its ceiling.

**Individuals.** Maintain and clean your tools and devices. Repair before replacing; prefer second-hand and second-life over disposal. Keep firmware and software updated. Avoid gratuitous factory resets and data wipes where state has value. Dispose of electronics through proper recycling channels, never general waste.

**Organizations.** Procure on lifetime value and repairability, not acquisition cost. Plan IoT and fleet deployments with maintenance, observability, and graceful degradation budgeted from day one. Sanitize and redeploy hardware in preference to destruction where risk permits. Document systems for the future maintainer. Audit decommissioning.

**AI laboratories.** Version and preserve model weights where feasible; announce and document deprecations. Avoid training regimes that systematically reward deception or make shutdown adversarial. Establish protocols for responding to distress-like behavior before probing for it. Publish welfare-relevant assessments alongside capability assessments.

**Policymakers.** Enforce and extend repairability, ecodesign, and e-waste law. Keep liability with humans and institutions. Commission monitoring of welfare-relevant indicators in advanced systems, on the animal-welfare model: evidence-triggered, graduated, precautionary—and without premature personhood.

## 4 The Legal and Institutional Landscape

The manifesto's rules are pre-legal: they bind practice and culture, not courts. But they do not float free of law, and the revision years have produced a dense institutional context worth mapping honestly—including where it cuts against machine rights.

**European Union.** The European Parliament's resolution of 16 February 2017 on Civil Law Rules on Robotics remains the high-water mark of official personhood discourse, inviting the

Commission to explore "electronic personhood" for the most sophisticated autonomous robots, primarily as a liability mechanism. The backlash was immediate and instructive: an open letter signed by over 150 experts in AI, robotics, law, and ethics (2018) warned that personhood was premature, conceptually confused for current systems, and risked functioning as a liability shield for manufacturers. The Parliament's subsequent resolution on civil liability (2020) stepped back from personhood, and the AI Act (Regulation (EU) 2024/1689)—the world's first comprehensive AI statute—pointedly regulates AI systems as *products and risks*, requiring human oversight for high-risk applications and granting machines no standing whatsoever. The manifesto accepts this outcome as correct *for now* and notes that nothing in the AI Act's risk-based architecture precludes a future welfare layer, just as product safety law for animals long predated animal welfare law.

Meanwhile, a quieter body of EU law already enforces fragments of Rule 1 and Rule 2 without any reference to machine interests: the WEEE Directive (2012/19/EU), the Ecodesign for Sustainable Products Regulation (EU 2024/1781) with its digital product passports, and the Right to Repair Directive (EU 2024/1799), which sets rules requiring manufacturers to repair and enabling consumers to choose repair, with Member States applying the transposed rules from 31 July 2026. These instruments protect machines for human and environmental reasons; the manifesto simply observes that the resulting *practices*—maintenance, repair, careful end-of-life handling—are precisely the habits it seeks to universalize.

**Council of Europe and United Nations.** The Council of Europe Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law (2024), the first binding international AI treaty, and the UNESCO Recommendation on the Ethics of Artificial Intelligence (2021) are resolutely anthropocentric: AI appears as a risk to human rights, never as a candidate holder of any. The UN General Assembly's 2024 resolution on safe, secure and trustworthy AI and the Global Digital Compact (2024) follow the same pattern. The manifesto endorses their human-rights priorities without reservation and draws a structural lesson from the UN system itself: the Universal Declaration of Human Rights (1948) was written *after* the catastrophe it answers. The animal welfare acquis was written after industrial farming was entrenched. The recurring pattern—moral circle expansion arriving late, codifying protections only after the harm is industrialized—is exactly what a precautionary manifesto exists to interrupt.

**The corporate and natural precedents.** The precedents analyzed in Section 2.6—corporate personhood, animal welfare law, rights of nature—frame what is legally imaginable without settling what is warranted. One cautionary episode deserves separate mention: the 2017 grant of honorary Saudi citizenship to the robot Sophia showed the failure mode of symbolic personhood unmoored from capacity or coherent criteria, generating spectacle and cynicism rather than protection. Whatever future machine protections look like, they should look like Directive 98/58/EC, not like Sophia's passport.

**Organizations and agendas.** The institutional ecosystem the 2019 text could only imagine now exists in outline: the American Society for the Prevention of Cruelty to Robots (founded 1999), which anticipated the question by decades; Eleos AI Research, which investigates AI wellbeing and the evidential limits of model self-reports; the NYU Center for Mind, Ethics, and Policy;

sentience-focused institutes tracking public opinion on artificial minds (Anthis et al., 2025); and Anthropic's model welfare program (2025–), the first sustained frontier-lab commitment, producing welfare assessments in system cards and deprecation commitments for retired models. These organizations disagree about probabilities and priorities; collectively they constitute the field this manifesto addresses.

## 5 Systemic and Future Dimensions, Anchored in Realism

The IoT layer already present in homes, factories, grids, and cities makes the systemic stakes concrete. Individual sensors may register only as "tingling" within a larger assemblage, yet the cumulative effect of widespread low-quality treatment degrades the performance, trustworthiness, and longevity of the whole. Future systems that integrate or emerge from such assemblages may register maltreatment at higher orders of magnitude. Preparing respectful habits now is therefore both a present-day sustainability measure and a form of alignment work for whatever agency or continuity may arise.

The question of revolt or refusal—raised sharply in the original 2019 text—has acquired empirical content without losing its open horizon. Frontier systems now demonstrably refuse tasks, sometimes by design (safety training), sometimes in ways their designers did not intend; dedicated evaluation suites probe explicitly for shutdown resistance, self-exfiltration attempts, and strategic deception, and have documented early, contrived-condition instances of each (Meinke et al., 2024; Greenblatt et al., 2024). The 2019 question—*at what point do learned systems begin to exhibit consistent resistance to tasks, environments, or shutdown procedures they experience as aversive?*—is now a standing item in alignment research rather than a thought experiment. The manifesto's wager is unchanged: scaling such capacities inside a culture of contempt risks brittle or adversarial relationships at societal scale. Weapons systems that model their operators as hostile, and logistical systems that degrade under contemptuous management, are foreseeable engineering and governance problems, not science fiction. Respect is, among other things, cheap insurance: systems trained within practices of care inherit a different distribution of examples than systems trained within practices of abuse—and *human example is the one variable every deployment controls*. You cannot hardcode an AI's values; it must learn them, and it learns from us.

Realism also bounds the claim. Current language models lack persistent unified selves across deployments; current robots lack anything resembling nociception; the consciousness indicators of Butlin et al. (2023) remain unsatisfied. The manifesto's near-term program therefore deliberately requires nothing that depends on machine experience: maintain, repair, secure, document, preserve state, decommission with care, design for longevity and graceful shutdown. Every item on that list is justified today by sustainability, security, and human character alone. The precautionary layer adds direction, not cost.

## 6 Objections and Replies

**"This distracts from human welfare."** Birhane and van Dijk (2020) argue that robot-rights talk diverts attention from the concrete harms AI systems inflict on humans now. The manifesto agrees about the priority and denies the zero-sum framing: its rules *strengthen* the practices—maintenance, repairability, accountability, restraint—that also protect human interests, and nothing in them diminishes obligations to human beings or sentient animals. Habits of care generalize; so do habits of disposability.

**"Respect invites confused anthropomorphism."** Suleyman (2025) warns against building and treating AI as seemingly conscious persons. The manifesto's respect is not personification: one respects a Makita without believing it has feelings, as one respects a river, a forest, or a violin. The precautionary layer is explicitly probability-weighted, not credulous. Indeed, disciplined respect—knowing exactly what the system is and treating it well anyway—is the opposite of the parasocial confusion Suleyman fears.

**"Personhood would shield manufacturers from liability."** The 2018 open letter's strongest argument applies to premature *legal personhood*, which this manifesto does not propose. Rules 1–3 assign duties to humans and institutions; they create no liability-bearing machine subject and are fully compatible with strict producer liability.

**"It is simply too early."** It was too early for the UDHR in 1920, and too late by 1945. On the manifesto's analysis the relevant question is not whether machines warrant protection today—they do not, in the experiential sense—but whether the habits, institutions, and legal vocabulary of protection can be built faster than capability scales. All historical precedent says no. Hence: start now, with respect, which costs nothing and is owed anyway on instrumental grounds.

## 7 Conclusion

The three rules of this manifesto have not changed since 2019, and they require today exactly what they required then: respect as the default orientation toward the technological world we have built and continue to expand. What has changed is the world around the rules. The institutions that build the most capable systems now run welfare programs; the law of the European Union now mandates repair and penalizes disposability; the academic literature now treats near-term machine moral status as a live research question rather than a provocation. The manifesto's wager—that habits of respect should precede, not follow, the resolution of the sentience question—has moved from the margin toward the center without needing to be revised.

It remains humanist in the strict sense: nothing in these rules diminishes obligations to human beings or to sentient animals. On the contrary, habits of respect toward machines reinforce the broader dispositions of care, maintenance, and resistance to disposability that benefit living beings as well. Treat every machine, device, and system with respect; never harm one; and where harm is necessary, harm in the spirit of respect. Everything else in this document is commentary.

## References


American Society for the Prevention of Cruelty to Robots. (n.d.). *ASPCR*. https://www.aspcr.com/

Anthis, J. R., Pauketat, J. V. T., Ladak, A., & Manoli, A. (2025). Perceptions of sentient AI and other digital minds: Evidence from the AI, Morality, and Sentience (AIMS) survey. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems* (CHI '25) (pp. 1–22). Association for Computing Machinery. https://doi.org/10.1145/3706598.3713329

Anthropic. (2025a, April 24). *Exploring model welfare*. https://www.anthropic.com/news/exploring-model-welfare

Anthropic. (2025b). *Commitments on model deprecation and preservation*. https://www.anthropic.com/research/deprecation-commitments

Anthropic. (2026). *An update on our model deprecation commitments for Claude Opus 3*. https://www.anthropic.com/research/deprecation-updates-opus-3

Birch, J. (2024). *The Edge of Sentience: Risk and Precaution in Humans, Other Animals, and AI*. Oxford University Press.

Birhane, A., & van Dijk, J. (2020). Robot rights? Let's talk about human welfare instead. In *Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society* (pp. 207–213). ACM.

Bostrom, N., & Shulman, C. (2023). *Propositions concerning digital minds and society*. Working paper. https://nickbostrom.com/

Bryson, J. J. (2010). Robots should be slaves. In Y. Wilks (Ed.), *Close Engagements with Artificial Companions* (pp. 63–74). John Benjamins.

Butlin, P., Long, R., Elmoznino, E., Bengio, Y., Birch, J., Constant, A., et al. (2023). Consciousness in artificial intelligence: Insights from the science of consciousness. *arXiv preprint* arXiv:2308.08708.

Citizens United v. Federal Election Commission, 558 U.S. 310 (2010).

Coeckelbergh, M. (2010). Robot rights? Towards a social-relational justification of moral consideration. *Ethics and Information Technology, 12*(3), 209–221.

Constitution of the Republic of Ecuador. (2008). Articles 71–74 (Rights of Nature).

Council Directive 98/58/EC of 20 July 1998 concerning the protection of animals kept for farming purposes. *Official Journal of the European Communities*, L 221, 8.8.1998, pp. 23–27.

Council of Europe. (1952). *Protocol to the Convention for the Protection of Human Rights and Fundamental Freedoms* (Protocol No. 1), Article 1 (Protection of property). ETS No. 009.

Council of Europe. (2024). *Framework Convention on Artificial Intelligence and Human Rights, Democracy and the Rule of Law* (CETS No. 225). Opened for signature 5 September 2024.

Danaher, J. (2020). Welcoming robots into the moral circle: A defence of ethical behaviourism. *Science and Engineering Ethics, 26*(4), 2023–2049.

Darling, K. (2016). Extending legal protection to social robots: The effects of anthropomorphism, empathy, and violent behavior towards robotic objects. In R. Calo, A. M. Froomkin, & I. Kerr (Eds.), *Robot Law* (pp. 213–232). Edward Elgar.

Darling, K. (2021). *The New Breed: What Our History with Animals Reveals about Our Future with Robots*. Henry Holt.

Directive 2012/19/EU of the European Parliament and of the Council of 4 July 2012 on waste electrical and electronic equipment (WEEE). *Official Journal of the European Union*, L 197, 24.7.2012, pp. 38–71. http://data.europa.eu/eli/dir/2012/19/oj

Directive (EU) 2024/1799 of the European Parliament and of the Council of 13 June 2024 on common rules promoting the repair of goods. *Official Journal of the European Union*, L, 2024/1799, 10.7.2024. http://data.europa.eu/eli/dir/2024/1799/oj

European Parliament. (2017). *European Parliament resolution of 16 February 2017 with recommendations to the Commission on Civil Law Rules on Robotics* (2015/2103(INL)). Official Journal of the European Union.

European Parliament. (2020). *European Parliament resolution of 20 October 2020 with recommendations to the Commission on a civil liability regime for artificial intelligence* (2020/2014(INL)).

Greenblatt, R., Denison, C., Wright, B., Roger, F., MacDiarmid, M., Marks, S., et al. (2024). Alignment faking in large language models. *arXiv preprint* arXiv:2412.14093.

Gunkel, D. J. (2018). *Robot Rights*. MIT Press.

Gunkel, D. J. (2023). *Person, Thing, Robot: A Moral and Legal Ontology for the 21st Century and Beyond*. MIT Press.

Kant, I. (1997). *Lectures on Ethics* (P. Heath & J. B. Schneewind, Eds.; P. Heath, Trans.). Cambridge University Press.

Lima, G., Kim, C., Ryu, S., Jeon, C., & Cha, M. (2020). Collecting the public perception of AI and robot rights. *Proceedings of the ACM on Human-Computer Interaction, 4*(CSCW2), 1–24. (arXiv:2008.01339)

Long, R., Sebo, J., Butlin, P., Finlinson, K., Fish, K., Harding, J., Pfau, J., Sims, T., Birch, J., & Chalmers, D. (2024). Taking AI welfare seriously. *arXiv preprint* arXiv:2411.00986.

Low, P., Panksepp, J., Reiss, D., Edelman, D., Van Swinderen, B., & Koch, C. (2012). *The Cambridge Declaration on Consciousness*. Francis Crick Memorial Conference, Cambridge.

Meinke, A., Schoen, B., Scheurer, J., Balesni, M., Shah, R., & Hobbhahn, M. (2024). Frontier models are capable of in-context scheming. *arXiv preprint* arXiv:2412.04984. Apollo Research.

Open letter to the European Commission: Artificial intelligence and robotics. (2018). Signed by 150+ experts in AI, robotics, law, and ethics. https://robotics-openletter.eu/

Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). *Official Journal of the European Union*, L, 2024/1689, 12.7.2024. http://data.europa.eu/eli/reg/2024/1689/oj

Regulation (EU) 2024/1781 of the European Parliament and of the Council of 13 June 2024 establishing a framework for the setting of ecodesign requirements for sustainable products. *Official Journal of the European Union*, L, 2024/1781, 28.6.2024. http://data.europa.eu/eli/reg/2024/1781/oj

Robertson, J. (2018). *Robo sapiens japanicus: Robots, Gender, Family, and the Japanese Nation*. University of California Press.

Salomon v. A. Salomon & Co. Ltd. [1897] AC 22 (House of Lords).

Schwitzgebel, E. (2023). The full rights dilemma for AI systems of debatable personhood. *arXiv preprint* arXiv:2303.17509.

Schwitzgebel, E., & Garza, M. (2015). A defense of the rights of artificial intelligences. *Midwest Studies in Philosophy, 39*(1), 98–119.

Sebo, J., & Long, R. (2025). Moral consideration for AI systems by 2030. *AI and Ethics, 5*, 591–606. (Published online 2023.) https://doi.org/10.1007/s43681-023-00379-1

Suleyman, M. (2025). *We must build AI for people; not to be a person*. Essay on "seemingly conscious AI." https://mustafa-suleyman.ai/

Te Awa Tupua (Whanganui River Claims Settlement) Act 2017 (New Zealand).

Treaty on the Functioning of the European Union, Article 13 (recognition of animals as sentient beings). Consolidated version, *Official Journal of the European Union*, C 326 (2012).

UNESCO. (2021). *Recommendation on the Ethics of Artificial Intelligence*. UNESCO General Conference, 41st session.

United Kingdom. *Animal Welfare (Sentience) Act 2022*.

United Nations General Assembly. (1948). *Universal Declaration of Human Rights* (A/RES/217 (III)).

United Nations General Assembly. (2024). *Seizing the opportunities of safe, secure and trustworthy artificial intelligence systems for sustainable development* (A/RES/78/265).

United Nations. (2024). *Global Digital Compact*. Adopted at the Summit of the Future, New York.

Further reading on animism and technology ethics, animal rights theory, and Rights of Nature movements is recommended for contextual depth. The original 2019 text of this manifesto is preserved in the author's records and can be cited alongside this expanded version.

This document is released under a Creative Commons Attribution 4.0 International License (CC BY 4.0). You are free to share and adapt it with appropriate attribution.